

PD1

Jiarong Ye, Yuan Meng

November 10, 2018

import packages

```
In [205]: import datascience as ds
          from datascience import *
          import numpy as np
          from graphviz import Source
          import pandas as pd
          import seaborn as sns
          from sklearn.pipeline import Pipeline
          from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
          from sklearn.tree import DecisionTreeClassifier
          from sklearn.ensemble import RandomForestClassifier
          from sklearn import tree
          from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score,
                                     accuracy_score, classification_report

          import matplotlib.pyplot as plt
          from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
          from sklearn.externals import joblib
          %matplotlib inline
```

tweets data loaded into Jupyter Notebook as Table object

```
In [26]: df = ds.Table.read_table('Climate1SupportiveLevel.csv', sep=',')
          df
```

```
Out[26]: Unnamed: 0 | ID | Text
0 | 962_Cleand_Climate1.csv | RT @kasseroles: Energy is the #1 contributor t
1 | 885_Cleand_Climate1.csv | RT @edelman_barbara: @msnbc why don t you have
2 | 680_Cleand_Climate1.csv | RT @OtagoGrad: @anthonyfurey @OskieOckham The d
3 | 1152_Cleand_Climate1.csv | The Dow just recorded its 3rd worst day ever. T
4 | 731_Cleand_Climate1.csv | RT @SimonBanksHB: I am not going to rule out th
5 | 1075_Cleand_Climate1.csv | RT @sydneyleemarco: nothing like an 80 degree o
6 | 85_Cleand_Climate1.csv | @MerlenesMemos @CNN It's not an act of god. Cli
7 | 654_Cleand_Climate1.csv | RT @MikeLevinCA: When asked about climate chang
8 | 916_Cleand_Climate1.csv | RT @gq_jayq: Bet I got 11 years to run it up ht
9 | 372_Cleand_Climate1.csv | No they care about the oil billionaires
... (1273 rows omitted)
```

Preprocess

```
In [276]: X = list(df['Text'])
          y = list(df['SupportiveLabel'])
```

Check whether the data distribution is balanced

```
In [89]: def check(sentiment, index, note='training'):
          if sentiment==0:
              label = 'not supportive'
          else:
              label = 'supportive'
          print('There are {} '.format(df.take(index).where('SupportiveLabel',
                                                             are.equal_to(sentiment)).size[0][0])+label+' tweets in the '+note+' set.')
```

Model Building

```
In [90]: def custom_split(train_index, test_index):
          trainingset = df.take(train_index)
          testingset = df.take(test_index)

          X_train= list(trainingset['Text'])
          y_train= list(trainingset['SupportiveLabel'])
          X_test= list(testingset['Text'])
          y_test= list(testingset['SupportiveLabel'])

          return X_train, X_test, y_train, y_test
```

classifier

```
In [291]: def classifier(X_train, y_train, X_test, fold, max_depth, min_samples_leaf):
          # token_pattern='(([#@]| [0-9]| [a-z]| [A-Z]))+'
          clf = Pipeline(
              [
                  ('vect', CountVectorizer(token_pattern="(!RT|rt|\\d+) [#]*[\\w\\'_-]{2,100}"
                                           analyzer = 'word',
                                           stop_words='english',
                                           min_df = 3)),
                  ('clf', DecisionTreeClassifier(criterion='entropy',
                                                  random_state = 100,
                                                  max_depth = max_depth,
                                                  min_samples_leaf = min_samples_leaf))
              ])
          clf.fit(X_train, y_train)
          feature_names = clf.named_steps['vect'].get_feature_names()
          try:
              dot_data = tree.export_graphviz(clf.named_steps['clf'], out_file=None,
                                              feature_names=feature_names)
              graph = Source(dot_data)
```

```

        graph.render('ClimateClassifier-Fold_{}'.format(fold))
    except Exception as e:
        print(e)
    predicted_y_train = clf.predict(X_train)
    predicted_y_test = clf.predict(X_test)
    # save as pickle
    joblib.dump(clf, 'ClimateTeam7PD1.pkl')
    return predicted_y_train, predicted_y_test

```

```

In [283]: c=CountVectorizer(token_pattern="(?!RT|rt|\d+)[@#]*[\w\ ' _-]{2,100}",
                             analyzer = 'word',
                             stop_words='english',
                             min_df = 3)

c.fit(X, y)
c.get_feature_names()

```

```

Out[283]: ['#1o5c',
           '#actonclimate',
           '#auspol',
           '#cdnpoli',
           '#climate',
           '#climateaction',
           '#climatebreakdown',
           '#climatechange',
           '#climatechangeisreal',
           '#climateimpactsvic',
           '#dems',
           '#emissions',
           '#energy',
           '#environment',
           '#florida',
           ....

```

evaluation

```

In [295]: def eval_results(predicted_y_train, y_train, predicted_y_test, y_test):
    accuracy_s = accuracy_score(y_test, predicted_y_test)
    precision_s = precision_score(y_test, predicted_y_test)
    recall_s = recall_score(y_test, predicted_y_test)
    f1_s = f1_score(y_test, predicted_y_test)
    cm_train = confusion_matrix(y_train, predicted_y_train)
    cm_test = confusion_matrix(y_test, predicted_y_test)

    print('Accuracy Score:', accuracy_s)
    print("Precision Score:", precision_s)
    print("Recall Score:", recall_s)
    print("f1 Score:", f1_s)
    print('confusion_matrix of training set is: \n', cm_train, '\n')

```

```

print('confusion_matrix of testing set is: \n', cm_test, '\n')
print(classification_report(y_test, predicted_y_test))

classes = ['not supportive', 'supportive']
sns.heatmap(cm_train, annot=True, cmap='Blues', yticklabels=classes,
            xticklabels=classes)

plt.title('confusion matrix of training set')
plt.show()
sns.heatmap(cm_test, annot=True, cmap='Blues', yticklabels=classes,
            xticklabels=classes)

plt.title('confusion matrix of testing set')
plt.show()
return accuracy_s, precision_s, recall_s, f1_s

```

k-fold

```

In [293]: def k_fold_evaluate(X, y, max_depth, min_samples_leaf):
    # initialization
    accuracy = []
    precision = []
    recall = []
    f1 = []
    fold = 1
    skf = StratifiedKFold(n_splits=5, random_state=1, shuffle= True)

    # build model and collect results
    for train_index, test_index in skf.split(X, y):
        if fold==1:
            list(map(lambda x: check(x, train_index), range(2)))
            list(map(lambda x: check(x, test_index, note='testing'), range(2)))

        X_train, X_test, y_train, y_test = custom_split(train_index, test_index)

        predicted_y_train, predicted_y_test =
                                                    classifier(X_train=X_train, y_train=y_train,
                                                                X_test=X_test, fold=fold,
                                                                max_depth = max_depth,
                                                                min_samples_leaf = min_samples_leaf)

        print('\nFold: {}'.format(fold))
        accuracy_s, precision_s, recall_s, f1_s = eval_results(predicted_y_train,
                                                                y_train, predicted_y_test, y_train)

        accuracy.append(accuracy_s)
        precision.append(precision_s)
        recall.append(recall_s)
        f1.append(f1_s)

```

```

metrics_df = pd.DataFrame(
    {
        'accuracy': accuracy,
        'precision': precision,
        'recall': recall,
        'f1': f1
    }
)
fold += 1
return metrics_df

```

Tasks

- A description of model parameters you tried and the associated Stratified k-fold cross validation results for each model parameter choice
- Describe the model parameters you chose and the rationale of your decision.
- Double check overfitting risk: Compare the model's confusion matrix for training data vs the model's confusion matrix for testing data.

In [296]: `k_fold_evaluate(X, y, max_depth=5, min_samples_leaf=2)`

There are 453 not supportive tweets in the training set.
 There are 572 supportive tweets in the training set.
 There are 114 not supportive tweets in the testing set.
 There are 144 supportive tweets in the testing set.

Fold: 1

Accuracy Score: 0.6472868217054264

Precision Score: 0.6853146853146853

Recall Score: 0.6805555555555556

f1 Score: 0.6829268292682927

confusion_matrix of training set is:

```
[[268 185]
```

```
[153 419]]
```



confusion_matrix of testing set is:

```
[[69 45]
```

```
[46 98]]
```

	precision	recall	f1-score	support
0	0.60	0.61	0.60	114
1	0.69	0.68	0.68	144
micro avg	0.65	0.65	0.65	258

macro avg	0.64	0.64	0.64	258
weighted avg	0.65	0.65	0.65	258





Fold: 2
 Accuracy Score: 0.6809338521400778
 Precision Score: 0.7046979865771812

Recall Score: 0.7342657342657343
 f1 Score: 0.7191780821917809
 confusion_matrix of training set is:
 [[263 190]
 [137 436]]

confusion_matrix of testing set is:
 [[70 44]
 [38 105]]

	precision	recall	f1-score	support
0	0.65	0.61	0.63	114
1	0.70	0.73	0.72	143
micro avg	0.68	0.68	0.68	257
macro avg	0.68	0.67	0.67	257
weighted avg	0.68	0.68	0.68	257





Fold: 3

Accuracy Score: 0.625

Precision Score: 0.6516129032258065

Recall Score: 0.7062937062937062

f1 Score: 0.6778523489932886

confusion_matrix of training set is:

```
[[273 181]
```



```
[135 438]]
```

confusion_matrix of testing set is:

```
[[ 59  54]
```

```
[ 42 101]]
```

	precision	recall	f1-score	support
0	0.58	0.52	0.55	113
1	0.65	0.71	0.68	143
micro avg	0.62	0.62	0.62	256
macro avg	0.62	0.61	0.61	256
weighted avg	0.62	0.62	0.62	256




Fold: 4
 Accuracy Score: 0.6015625
 Precision Score: 0.6289308176100629
 Recall Score: 0.6993006993006993
 f1 Score: 0.6622516556291391
 confusion_matrix of training set is:
 [[275 179]
 [136 437]]


confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 43 100]]
```

	precision	recall	f1-score	support
0	0.56	0.48	0.51	113
1	0.63	0.70	0.66	143
micro avg	0.60	0.60	0.60	256
macro avg	0.59	0.59	0.59	256
weighted avg	0.60	0.60	0.60	256





Fold: 5

Accuracy Score: 0.63671875

Precision Score: 0.6488095238095238

Recall Score: 0.7622377622377622

f1 Score: 0.7009646302250804

confusion_matrix of training set is:

```
[[272 182]
```



```
[137 436]]
```

confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 34 109]]
```

	precision	recall	f1-score	support
0	0.61	0.48	0.54	113
1	0.65	0.76	0.70	143
micro avg	0.64	0.64	0.64	256
macro avg	0.63	0.62	0.62	256
weighted avg	0.63	0.64	0.63	256



```
Out [296]:
```

	accuracy	precision	recall	f1
0	0.647287	0.685315	0.680556	0.682927
1	0.680934	0.704698	0.734266	0.719178
2	0.625000	0.651613	0.706294	0.677852
3	0.601562	0.628931	0.699301	0.662252
4	0.636719	0.648810	0.762238	0.700965

```
In [297]: k_fold_evaluate(X, y, max_depth=6, min_samples_leaf=2)
```

There are 453 not supportive tweets in the training set.
 There are 572 supportive tweets in the training set.
 There are 114 not supportive tweets in the testing set.
 There are 144 supportive tweets in the testing set.

Fold: 1

Accuracy Score: 0.6434108527131783

Precision Score: 0.678082191780822

Recall Score: 0.6875

f1 Score: 0.6827586206896552

confusion_matrix of training set is:

[[273 180]


[150 422]]


confusion_matrix of testing set is:

[[67 47]

[45 99]]

	precision	recall	f1-score	support
0	0.60	0.59	0.59	114
1	0.68	0.69	0.68	144
micro avg	0.64	0.64	0.64	258
macro avg	0.64	0.64	0.64	258
weighted avg	0.64	0.64	0.64	258





Fold: 2

Accuracy Score: 0.6770428015564203

Precision Score: 0.7027027027027027

Recall Score: 0.7272727272727273

f1 Score: 0.7147766323024054

confusion_matrix of training set is:

```
[[267 186]
```



```
[133 440]]
```

confusion_matrix of testing set is:

```
[[ 70  44]
```

```
[ 39 104]]
```

	precision	recall	f1-score	support
0	0.64	0.61	0.63	114
1	0.70	0.73	0.71	143
micro avg	0.68	0.68	0.68	257
macro avg	0.67	0.67	0.67	257
weighted avg	0.68	0.68	0.68	257




Fold: 3
 Accuracy Score: 0.61328125
 Precision Score: 0.6428571428571429
 Recall Score: 0.6923076923076923
 f1 Score: 0.6666666666666666
 confusion_matrix of training set is:
 [[280 174]
 [132 441]]


confusion_matrix of testing set is:

```
[[58 55]
```

```
[44 99]]
```

	precision	recall	f1-score	support
0	0.57	0.51	0.54	113
1	0.64	0.69	0.67	143
micro avg	0.61	0.61	0.61	256
macro avg	0.61	0.60	0.60	256
weighted avg	0.61	0.61	0.61	256





Fold: 4

Accuracy Score: 0.609375

Precision Score: 0.6335403726708074

Recall Score: 0.7132867132867133

f1 Score: 0.6710526315789473

confusion_matrix of training set is:

```
[[272 182]
```



```
[125 448]]
```

confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 41 102]]
```

	precision	recall	f1-score	support
0	0.57	0.48	0.52	113
1	0.63	0.71	0.67	143
micro avg	0.61	0.61	0.61	256
macro avg	0.60	0.60	0.60	256
weighted avg	0.60	0.61	0.60	256




Fold: 5
 Accuracy Score: 0.63671875
 Precision Score: 0.6488095238095238
 Recall Score: 0.7622377622377622
 f1 Score: 0.7009646302250804
 confusion_matrix of training set is:
 [[278 176]
 [133 440]]


confusion_matrix of testing set is:

```
[[ 54  59]
```

```
 [ 34 109]]
```

	precision	recall	f1-score	support
0	0.61	0.48	0.54	113
1	0.65	0.76	0.70	143
micro avg	0.64	0.64	0.64	256
macro avg	0.63	0.62	0.62	256
weighted avg	0.63	0.64	0.63	256





```
Out[297]:
```

	accuracy	precision	recall	f1
0	0.643411	0.678082	0.687500	0.682759
1	0.677043	0.702703	0.727273	0.714777
2	0.613281	0.642857	0.692308	0.666667
3	0.609375	0.633540	0.713287	0.671053
4	0.636719	0.648810	0.762238	0.700965

```
In [298]: k_fold_evaluate(X, y, max_depth=7, min_samples_leaf=1)
```

There are 453 not supportive tweets in the training set.
 There are 572 supportive tweets in the training set.
 There are 114 not supportive tweets in the testing set.
 There are 144 supportive tweets in the testing set.

Fold: 1

Accuracy Score: 0.6434108527131783

Precision Score: 0.678082191780822

Recall Score: 0.6875

f1 Score: 0.6827586206896552

confusion_matrix of training set is:

```
[[276 177]
```

```
[136 436]]
```


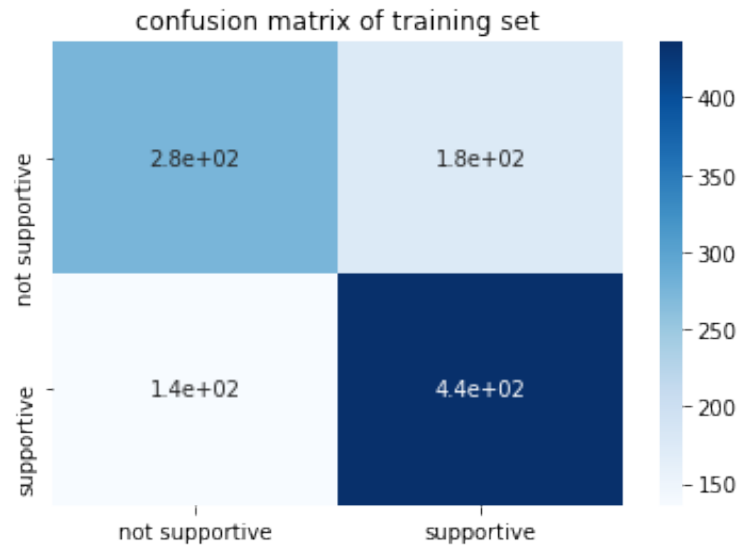
confusion_matrix of testing set is:

```
[[67 47]
```

```
[45 99]]
```

precision	recall	f1-score	support
-----------	--------	----------	---------


	0	0.60	0.59	0.59	114
	1	0.68	0.69	0.68	144
micro avg		0.64	0.64	0.64	258
macro avg		0.64	0.64	0.64	258
weighted avg		0.64	0.64	0.64	258




Fold: 2
 Accuracy Score: 0.6809338521400778
 Precision Score: 0.7103448275862069
 Recall Score: 0.7202797202797203
 f1 Score: 0.7152777777777778
 confusion_matrix of training set is:
 [[280 173]
 [139 434]]

confusion_matrix of testing set is:
 [[72 42]
 [40 103]]

	precision	recall	f1-score	support
0	0.64	0.63	0.64	114
1	0.71	0.72	0.72	143
micro avg	0.68	0.68	0.68	257
macro avg	0.68	0.68	0.68	257
weighted avg	0.68	0.68	0.68	257





Fold: 3

Accuracy Score: 0.61328125

Precision Score: 0.6447368421052632

Recall Score: 0.6853146853146853

f1 Score: 0.6644067796610169

confusion_matrix of training set is:

```
[[286 168]
```



```
[126 447]]
```

confusion_matrix of testing set is:

```
[[59 54]
```

```
[45 98]]
```

	precision	recall	f1-score	support
0	0.57	0.52	0.54	113
1	0.64	0.69	0.66	143
micro avg	0.61	0.61	0.61	256
macro avg	0.61	0.60	0.60	256
weighted avg	0.61	0.61	0.61	256




Fold: 4
 Accuracy Score: 0.60546875
 Precision Score: 0.63125
 Recall Score: 0.7062937062937062
 f1 Score: 0.6666666666666665
 confusion_matrix of training set is:
 [[283 171]
 [122 451]]


confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 42 101]]
```

	precision	recall	f1-score	support
0	0.56	0.48	0.52	113
1	0.63	0.71	0.67	143
micro avg	0.61	0.61	0.61	256
macro avg	0.60	0.59	0.59	256
weighted avg	0.60	0.61	0.60	256





Fold: 5

Accuracy Score: 0.6328125

Precision Score: 0.6467065868263473

Recall Score: 0.7552447552447552

f1 Score: 0.6967741935483871

confusion_matrix of training set is:

```
[[281 173]
```



```
[129 444]]
```

confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 35 108]]
```

	precision	recall	f1-score	support
0	0.61	0.48	0.53	113
1	0.65	0.76	0.70	143
micro avg	0.63	0.63	0.63	256
macro avg	0.63	0.62	0.62	256
weighted avg	0.63	0.63	0.63	256



```
Out [298]: accuracy precision recall f1
0 0.643411 0.678082 0.687500 0.682759
1 0.680934 0.710345 0.720280 0.715278
2 0.613281 0.644737 0.685315 0.664407
3 0.605469 0.631250 0.706294 0.666667
4 0.632812 0.646707 0.755245 0.696774
```

chosen parameters: max_depth=7, min_samples_leaf=2

Because the 5-fold f1 score of max_depth=7, min_samples_leaf=2 are better.

```
In [301]: k_fold_evaluate(X, y, max_depth=7, min_samples_leaf=2)
```

There are 453 not supportive tweets in the training set.

There are 572 supportive tweets in the training set.

There are 114 not supportive tweets in the testing set.

There are 144 supportive tweets in the testing set.

Fold: 1

Accuracy Score: 0.6550387596899225

Precision Score: 0.6870748299319728

Recall Score: 0.7013888888888888

f1 Score: 0.6941580756013744

confusion_matrix of training set is:

```
[[273 180]
```


```
[146 426]]
```


confusion_matrix of testing set is:

```
[[ 68  46]
```

```
[ 43 101]]
```

	precision	recall	f1-score	support
0	0.61	0.60	0.60	114
1	0.69	0.70	0.69	144
micro avg	0.66	0.66	0.66	258
macro avg	0.65	0.65	0.65	258
weighted avg	0.65	0.66	0.65	258





Fold: 2

Accuracy Score: 0.6770428015564203

Precision Score: 0.7054794520547946

Recall Score: 0.7202797202797203

f1 Score: 0.71280276816609

confusion_matrix of training set is:

```
[[278 175]
```



```
[139 434]]
```

confusion_matrix of testing set is:

```
[[ 71  43]
```

```
[ 40 103]]
```

	precision	recall	f1-score	support
0	0.64	0.62	0.63	114
1	0.71	0.72	0.71	143
micro avg	0.68	0.68	0.68	257
macro avg	0.67	0.67	0.67	257
weighted avg	0.68	0.68	0.68	257




Fold: 3
 Accuracy Score: 0.62109375
 Precision Score: 0.6513157894736842
 Recall Score: 0.6923076923076923
 f1 Score: 0.6711864406779661
 confusion_matrix of training set is:
 [[285 169]
 [128 445]]


confusion_matrix of testing set is:

```
[[60 53]
```

```
[44 99]]
```

	precision	recall	f1-score	support
0	0.58	0.53	0.55	113
1	0.65	0.69	0.67	143
micro avg	0.62	0.62	0.62	256
macro avg	0.61	0.61	0.61	256
weighted avg	0.62	0.62	0.62	256





Fold: 4

Accuracy Score: 0.60546875

Precision Score: 0.63125

Recall Score: 0.7062937062937062

f1 Score: 0.6666666666666665

confusion_matrix of training set is:

```
[[281 173]
```



```
[124 449]]
```

confusion_matrix of testing set is:

```
[[ 54  59]
```

```
[ 42 101]]
```

	precision	recall	f1-score	support
0	0.56	0.48	0.52	113
1	0.63	0.71	0.67	143
micro avg	0.61	0.61	0.61	256
macro avg	0.60	0.59	0.59	256
weighted avg	0.60	0.61	0.60	256




Fold: 5
 Accuracy Score: 0.6328125
 Precision Score: 0.6467065868263473
 Recall Score: 0.7552447552447552
 f1 Score: 0.6967741935483871
 confusion_matrix of training set is:
 [[280 174]
 [129 444]]


confusion_matrix of testing set is:

```
[[ 54  59]
```

```
 [ 35 108]]
```

	precision	recall	f1-score	support
0	0.61	0.48	0.53	113
1	0.65	0.76	0.70	143
micro avg	0.63	0.63	0.63	256
macro avg	0.63	0.62	0.62	256
weighted avg	0.63	0.63	0.63	256

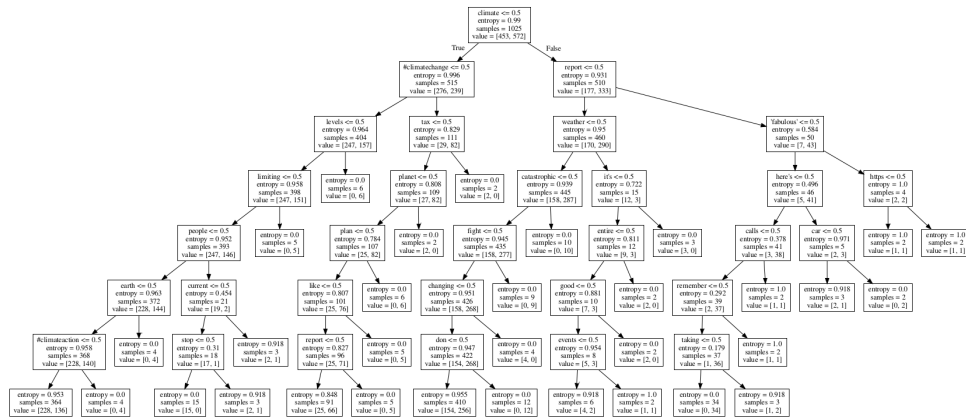




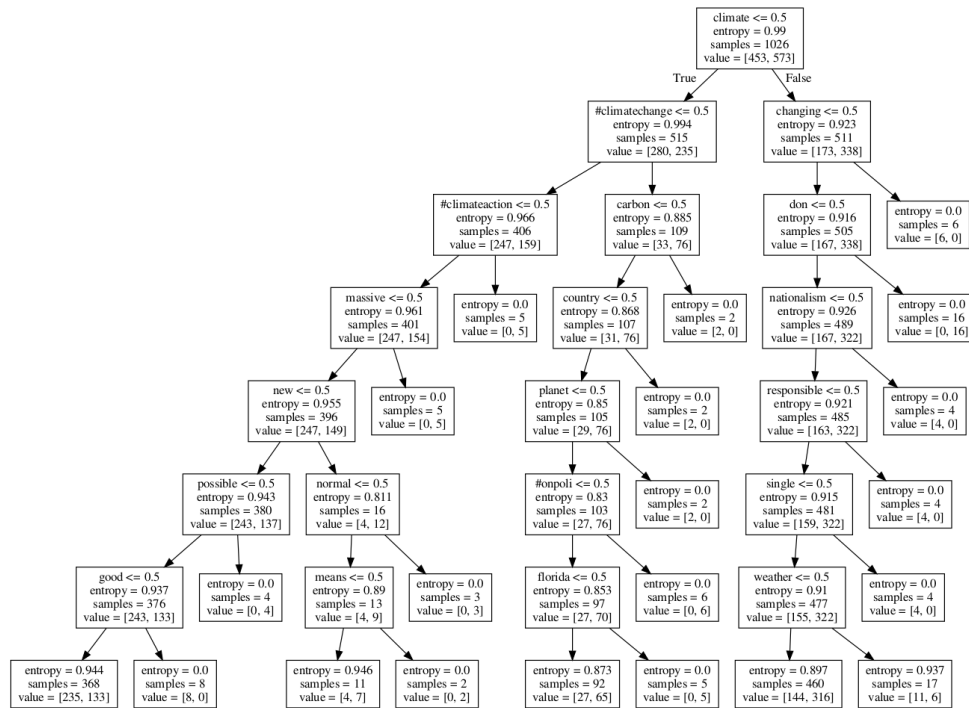
Out [301]:

	accuracy	precision	recall	f1
0	0.655039	0.687075	0.701389	0.694158
1	0.677043	0.705479	0.720280	0.712803
2	0.621094	0.651316	0.692308	0.671186
3	0.605469	0.631250	0.706294	0.666667
4	0.632812	0.646707	0.755245	0.696774

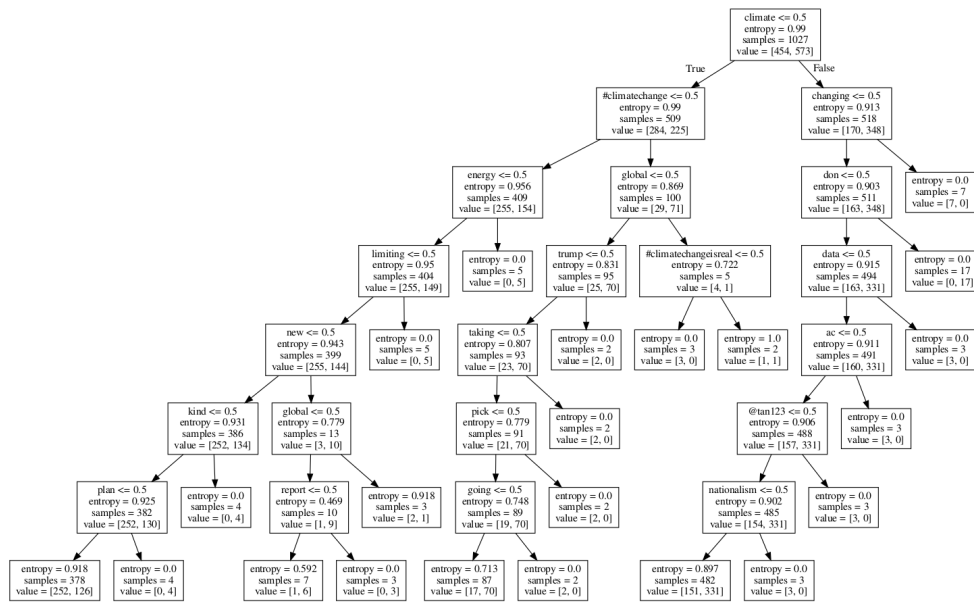
Fold 1



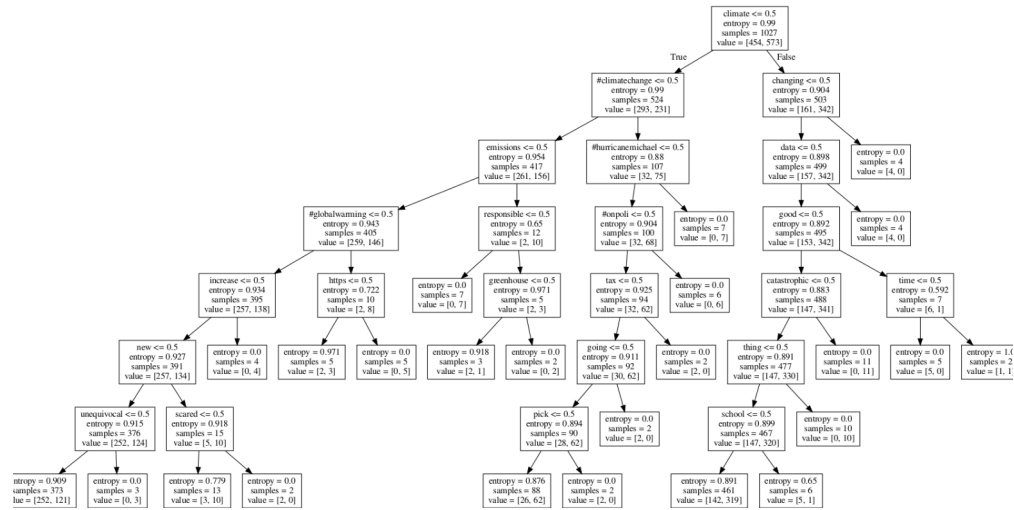
Fold 2



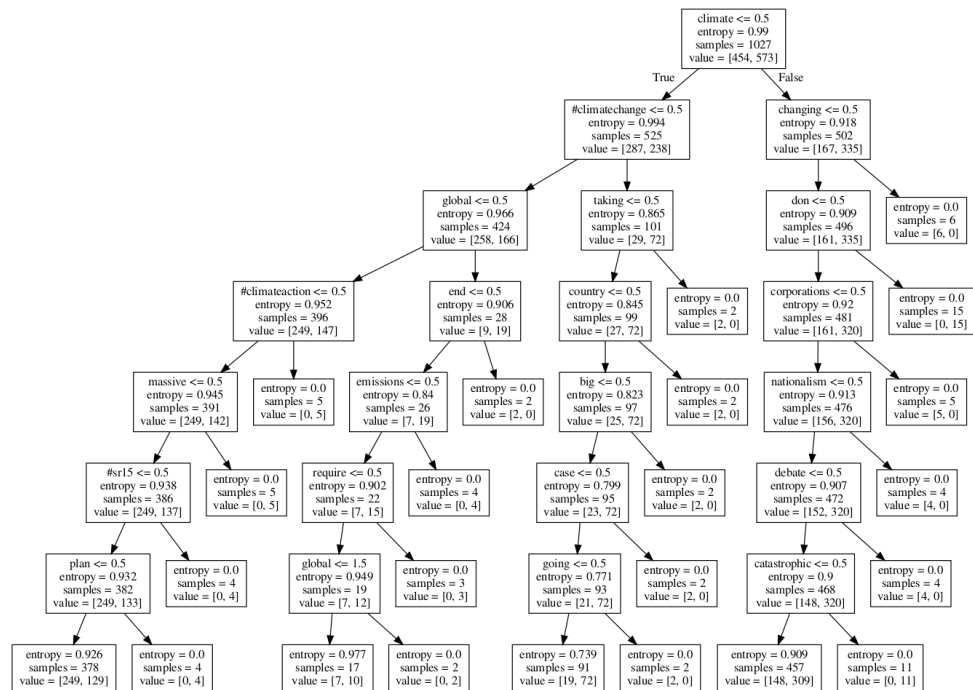
Fold 3



Fold 4



Fold 5



reload pickle

```
In [300]: clf2 = joblib.load('ClimateTeam7PD1.pkl')
          y_pred = clf2.predict(X)
          f1_score(y, y_pred)
```

```
Out[300]: 0.7355096602265155
```