

PD2

Team 7: Jiarong Ye, Yuan Meng

November 28, 2018

Project PD2

import packages

```
In [106]: import datascience as ds
          from datascience import *
          import numpy as np
          from collections import Counter
          from graphviz import Source
          import pandas as pd
          import seaborn as sns
          from sklearn.pipeline import Pipeline
          from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
          from sklearn.tree import DecisionTreeClassifier
          from sklearn.ensemble import RandomForestClassifier
          from sklearn import tree
          from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, accuracy_score
          import matplotlib.pyplot as plt
          from sklearn.model_selection import train_test_split, cross_val_score, StratifiedKFold
          from sklearn.externals import joblib
          %matplotlib inline
```

tweets data loaded into Jupyter Notebook as Table object

```
In [107]: df1 = ds.Table.read_table('Climate1SupportiveLevel.csv', sep=',')
          df2 = ds.Table.read_table('ClimateBalancedDS2.csv', sep=',')
          df = df1.append(df2)
          test_index = np.random.choice(df.num_rows, 100, replace=False)
          train_val_index = [i for i in np.arange(df.num_rows) if i not in test_index]
          test_data = df.take[test_index]
          df = df.take[train_val_index]
          X = list(df['Text'])
          y = list(df['Support'])
```

A description of all model enhancements incorporated into the construction of PD2.

used parameters

- CountVector

- analyzer = 'word'
- stop_words = the top 30 most common words in DS1+DS2 tweets
- min_df = 3
- ngram_range=(1,2)

- DecisionTree

- criterion='entropy'
- max_depth =
- min_samples_leaf = 2

A description of all model parameters you tried and the associated Stratified k-fold cross validation results for each model parameter choice based on the combined data set (DS1 + DS2)

- max_depth=range(1, 15)
- stop_words:
 - default in CountVectorizer: 'english'
 - user_defined: ['a', 'an', 'the', 'it', 'is', 'are', 'be', 'of', 'this', 'that', 'RT', 'rt', 'https']
 - the top 30 most common words in DS1+DS2 tweets (chosen)

code and test results as follows:

Check whether the data distribution is balanced

```
In [108]: def check(sentiment, index, note='training'):
            if sentiment==0:
                label = 'not supportive'
            else:
                label = 'supportive'
            print('There are {} '.format(df.take(index).where('Support',
                are.equal_to(sentiment)).size[0][0])+label+' tweets in the '+note+' set.')
```

Model Building

```
In [109]: def custom_split(train_index, test_index):
            trainingset = df.take(train_index)
            testingset = df.take(test_index)

            X_train= list(trainingset['Text'])
            y_train= list(trainingset['Support'])
            X_test= list(testingset['Text'])
            y_test= list(testingset['Support'])

            return X_train, X_test, y_train, y_test
```

classifier

```
In [110]: def classifier(X_train, y_train, X_test, fold, max_depth, min_samples_leaf, stop_words, overfit_risk):
# token_pattern='(([#@]|/[0-9]/[a-z]/[A-Z]))+'
clf = Pipeline(
    [
        ('vect', CountVectorizer(token_pattern="(!RT|rt|\\d+)[@#]*[\\w\\'\\_\\-]{2,100}",
                                analyzer = 'word',
                                stop_words = stop_words,
                                min_df = 3,
                                ngram_range=(1,2))),
        ('clf', DecisionTreeClassifier(criterion='entropy',
                                       random_state = 100,
                                       max_depth = max_depth,
                                       min_samples_leaf = min_samples_leaf))
    ])
clf.fit(X_train, y_train)
feature_names = clf.named_steps['vect'].get_feature_names()
try:
    dot_data = tree.export_graphviz(clf.named_steps['clf'], out_file=None,
                                    feature_names=feature_names)

    graph = Source(dot_data)
    graph.render('ClimateClassifier-Fold_{}'.format(fold))
except Exception as e:
    print(e)
predicted_y_train = clf.predict(X_train)
predicted_y_test = clf.predict(X_test)
# save as pickle
if overfit_risk:
    joblib.dump(clf, 'ClimateTeam7PD2_maxdepth{}.pkl'.format(max_depth))
else:
    joblib.dump(clf, 'ClimateTeam7PD2.pkl')
return predicted_y_train, predicted_y_test
```

evaluation

```
In [111]: def eval_results(predicted_y_train, y_train, predicted_y_test, y_test):
accuracy_s = accuracy_score(y_test, predicted_y_test)
precision_s = precision_score(y_test, predicted_y_test)
recall_s = recall_score(y_test, predicted_y_test)
f1_s = f1_score(y_test, predicted_y_test)
cm_train = confusion_matrix(y_train, predicted_y_train)
cm_test = confusion_matrix(y_test, predicted_y_test)

print('Accuracy Score:', accuracy_s)
print("Precision Score:", precision_s)
print("Recall Score:", recall_s)
print("f1 Score:", f1_s)
```

```

print('confusion_matrix of training set is: \n', cm_train, '\n')
print('confusion_matrix of testing set is: \n', cm_test, '\n')
print(classification_report(y_test, predicted_y_test))

classes = ['not supportive', 'supportive']
sns.heatmap(cm_train, annot=True, cmap='Blues', yticklabels=classes,
            xticklabels=classes)
plt.title('confusion matrix of training set')
plt.show()
sns.heatmap(cm_test, annot=True, cmap='Blues', yticklabels=classes,
            xticklabels=classes)
plt.title('confusion matrix of testing set')
plt.show()
return accuracy_s, precision_s, recall_s, f1_s

```

k-fold

```

In [112]: def k_fold_evaluate(X, y, max_depth, min_samples_leaf, stop_words, print_eval=True,
                             overfit_risk=False):
    # initialization
    accuracy = []
    precision = []
    recall=[]
    f1 = []
    fold = 1
    skf = StratifiedKFold(n_splits=5, random_state=1, shuffle= True)

    # build model and collect results
    for val_index, test_index in skf.split(X, y):
        X_train, X_val, y_train, y_val = custom_split(val_index, test_index)

        predicted_y_train, predicted_y_val = classifier(X_train=X_train, y_train=y_train,
                                                         X_test=X_val, fold=fold,
                                                         max_depth = max_depth,
                                                         min_samples_leaf = min_samples_leaf,
                                                         stop_words = stop_words,
                                                         overfit_risk=overfit_risk)

        metrics_df={}
        if print_eval:
            print('\nFold: {}'.format(fold))
            accuracy_s, precision_s, recall_s, f1_s = eval_results(predicted_y_train,
                                                                    y_train, predicted_y_val, y_val)

            accuracy.append(accuracy_s)
            precision.append(precision_s)
            recall.append(recall_s)
            f1.append(f1_s)

        metrics_df = pd.DataFrame(

```

```

        {
            'accuracy': accuracy,
            'precision': precision,
            'recall': recall,
            'f1': f1
        }
    )
    fold += 1
    return metrics_df

```

Tests:

test1: stop-words: default

```

In [113]: test_X = list(test_data['Text'])
          test_y = list(test_data['Support'])

```

```

In [117]: f1_lst_test1 = []
          f1_lst_train1 = []
          for d in range(1, 15):
              k_fold_evaluate(X, y, max_depth=d, min_samples_leaf=2, stop_words='english', print_eval
              clf_tmp = joblib.load('ClimateTeam7PD2_maxdepth{}.pkl'.format(d))
              print('maxdepth=', d)
              # test
              y_pred = clf_tmp.predict(test_X)
              print('test f1')
              print(f1_score(y_pred=y_pred, y_true=test_y))
              f1_lst_test1.append(f1_score(y_pred=y_pred, y_true=test_y))
              # train_val
              y_pred = clf_tmp.predict(X)
              print('train f1')
              print(f1_score(y_pred=y_pred, y_true=y))
              f1_lst_train1.append(f1_score(y_pred=y_pred, y_true=y))

```

```

maxdepth= 1
test f1
0.6019417475728155
train f1
0.6012009890498058
maxdepth= 2
test f1
0.6608695652173913
train f1
0.6829733163913596
maxdepth= 3
test f1
0.6666666666666666
train f1
0.6910466582597731

```

```
maxdepth= 4
test f1
0.6666666666666666
train f1
0.6962822936357909
maxdepth= 5
test f1
0.6782608695652174
train f1
0.7033312382149591
maxdepth= 6
test f1
0.6842105263157895
train f1
0.7010834926704909
maxdepth= 7
test f1
0.6956521739130435
train f1
0.7089410272669627
maxdepth= 8
test f1
0.6956521739130435
train f1
0.7140600315955766
maxdepth= 9
test f1
0.6956521739130435
train f1
0.7180296810862015
maxdepth= 10
test f1
0.6842105263157895
train f1
0.7214669617451785
maxdepth= 11
test f1
0.6956521739130435
train f1
0.7292057535959976
maxdepth= 12
test f1
0.6956521739130435
train f1
0.7314465408805032
maxdepth= 13
test f1
0.6956521739130435
```

```

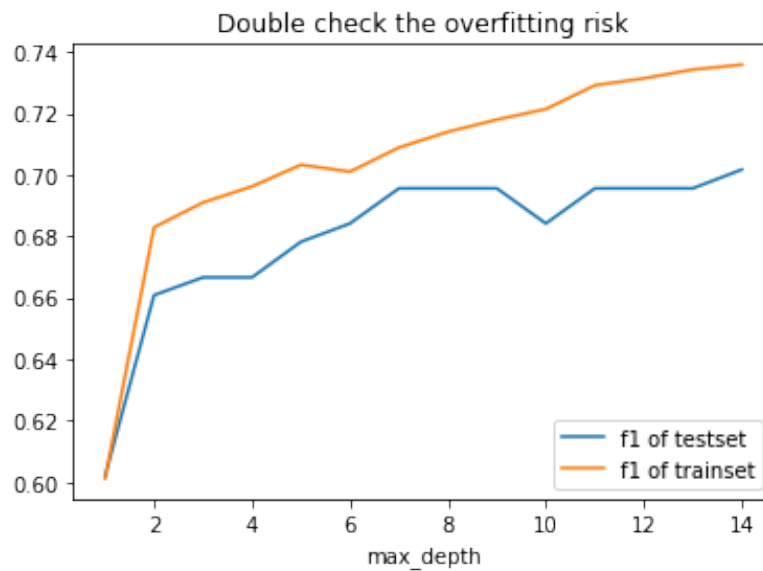
train f1
0.7343208320201702
maxdepth= 14
test f1
0.7017543859649122
train f1
0.7359143846395971

```

```

In [119]: plt.plot(np.arange(1,15), f1_lst_test1, label='f1 of testset')
plt.plot(np.arange(1,15), f1_lst_train1, label='f1 of trainset')
plt.xlabel('max_depth')
plt.title('Double check the overfitting risk')
plt.legend(loc='lower right')
plt.show()

```



test2: stop_words: ['a', 'an', 'the', 'it', 'is', 'are', 'be', 'of', 'this', 'that', 'RT', 'rt', 'https']

```

In [121]: f1_lst_test2 = []
f1_lst_train2 = []
for d in range(1, 15):
    k_fold_evaluate(X, y, max_depth=d, min_samples_leaf=2,
                    stop_words=['a', 'an', 'the', 'it', 'is', 'are', 'be', 'of', 'this', 't
                    print_eval=False, overfit_risk=True)
    clf_tmp = joblib.load('ClimateTeam7PD2_maxdepth{}.pkl'.format(d))
    print('maxdepth=', d)
    # test
    y_pred = clf_tmp.predict(test_X)

```

```

print('test f1')
print(f1_score(y_pred=y_pred, y_true=test_y))
f1_lst_test2.append(f1_score(y_pred=y_pred, y_true=test_y))
# train_val
y_pred = clf_tmp.predict(X)
print('train f1')
print(f1_score(y_pred=y_pred, y_true=y))
f1_lst_train2.append(f1_score(y_pred=y_pred, y_true=y))

```

```

maxdepth= 1
test f1
0.6019417475728155
train f1
0.6012009890498058
maxdepth= 2
test f1
0.6608695652173913
train f1
0.6829733163913596
maxdepth= 3
test f1
0.6666666666666665
train f1
0.7001540832049308
maxdepth= 4
test f1
0.6666666666666665
train f1
0.7059902200488998
maxdepth= 5
test f1
0.6724137931034483
train f1
0.704040087691826
maxdepth= 6
test f1
0.6608695652173913
train f1
0.704713049054184
maxdepth= 7
test f1
0.6608695652173913
train f1
0.7053915275994865
maxdepth= 8
test f1
0.6724137931034483
train f1

```



```


0.7151819322459222
maxdepth= 9
test f1
0.6724137931034483
train f1
0.7177522349936144
maxdepth= 10
test f1
0.6724137931034483
train f1
0.727102219443576
maxdepth= 11
test f1
0.6608695652173913
train f1
0.7317839195979898
maxdepth= 12
test f1
0.6551724137931034
train f1
0.74280408542247
maxdepth= 13
test f1
0.6666666666666665
train f1
0.747896540978498
maxdepth= 14
test f1
0.6608695652173913
train f1
0.7457735247208931

```

```

In [122]: plt.plot(np.arange(1,15), f1_lst_test2, label='f1 of testset')
          plt.plot(np.arange(1,15), f1_lst_train2, label='f1 of trainset')
          plt.xlabel('max_depth')
          plt.title('Double check the overfitting risk')
          plt.legend(loc='lower right')
          plt.show()

```



test3: stop_words: most common words in DS1+DS2 tweets

```
In [123]: stop_w = [i[0] for i in Counter([word for sentence in X for word in sentence.split()
                                          if 'climate' not in word.lower()
                                          and word.isalpha()
                                          and len(word)>1]).most_common():30]]
```

stop_w

```
Out[123]: ['the',
            'to',
            'RT',
            'of',
            'change',
            'is',
            'and',
            'in',
            'that',
            'on',
            'for',
            'are',
            'you',
            'we',
            'The',
            'it',
            'this',
            'about',
            'be',
            'by',
```

```

'have',
'not',
'will',
'our',
'from',
'as',
'can',
'with',
'all',
'We']

```

```

In [124]: f1_lst_test3 = []
          f1_lst_train3 = []
          for d in range(1, 15):
              k_fold_evaluate(X, y, max_depth=d, min_samples_leaf=2,
                              stop_words=stop_w,
                              print_eval=False, overfit_risk=True)
              clf_tmp = joblib.load('ClimateTeam7PD2_maxdepth{}.pkl'.format(d))
              print('maxdepth=', d)
              # test
              y_pred = clf_tmp.predict(test_X)
              print('test f1')
              print(f1_score(y_pred=y_pred, y_true=test_y))
              f1_lst_test3.append(f1_score(y_pred=y_pred, y_true=test_y))
              # train_val
              y_pred = clf_tmp.predict(X)
              print('train f1')
              print(f1_score(y_pred=y_pred, y_true=y))
              f1_lst_train3.append(f1_score(y_pred=y_pred, y_true=y))

```

```


maxdepth= 1
test f1
0.6019417475728155
train f1
0.6012009890498058
maxdepth= 2
test f1
0.6608695652173913
train f1
0.6829733163913596
maxdepth= 3
test f1
0.6666666666666666
train f1
0.6925750394944707
maxdepth= 4
test f1
0.6666666666666666

```

```
train f1
0.6969124133585382
maxdepth= 5
test f1
0.6724137931034483
train f1
0.70282131661442
maxdepth= 6
test f1
0.6724137931034483
train f1
0.7016661427224143
maxdepth= 7
test f1
0.6896551724137931
train f1
0.7094954559699153
maxdepth= 8
test f1
0.6896551724137931
train f1
0.7131249999999999
maxdepth= 9
test f1
0.6896551724137931
train f1
0.716875
maxdepth= 10
test f1
0.6896551724137931
train f1
0.7177722152690864
maxdepth= 11
test f1
0.7008547008547009
train f1
0.7199248120300753
maxdepth= 12
test f1
0.6896551724137931
train f1
0.7233238904627006
maxdepth= 13
test f1
0.7008547008547009
train f1
0.725434439178515
maxdepth= 14
```

```
test f1
0.706896551724138
train f1
0.7292782855341946
```

```
In [125]: plt.plot(np.arange(1,15), f1_lst_test3, label='f1 of testset')
plt.plot(np.arange(1,15), f1_lst_train3, label='f1 of trainset')
plt.xlabel('max_depth')
plt.title('Double check the overfitting risk')
plt.legend(loc='lower right')
plt.show()
```



Final result

```
In [126]: k_fold_evaluate(X, y, max_depth=15, min_samples_leaf=2,
                        stop_words=stop_w,
                        print_eval=True, overfit_risk=False)
```


```
Fold: 1
Accuracy Score: 0.6229802513464991
Precision Score: 0.6175637393767706
Recall Score: 0.7440273037542662
f1 Score: 0.6749226006191951
confusion_matrix of training set is:
[[649 407]
 [209 961]]
```


confusion_matrix of testing set is:

```
[[129 135]
```

```
[ 75 218]]
```

	precision	recall	f1-score	support
0	0.63	0.49	0.55	264
1	0.62	0.74	0.67	293
avg / total	0.62	0.62	0.62	557





Fold: 2

Accuracy Score: 0.6157989228007181

Precision Score: 0.6064690026954178

Recall Score: 0.7679180887372014

f1 Score: 0.6777108433734939

confusion_matrix of training set is:

```
[[619 437]
```



```
[200 970]]
```

confusion_matrix of testing set is:

```
[[118 146]
```

```
[ 68 225]]
```

	precision	recall	f1-score	support
0	0.63	0.45	0.52	264
1	0.61	0.77	0.68	293
avg / total	0.62	0.62	0.61	557



Fold: 3

Accuracy Score: 0.6481149012567325

Precision Score: 0.6456456456456456

Recall Score: 0.7337883959044369

f1 Score: 0.6869009584664537


confusion_matrix of training set is:



```
[[649 407]
 [220 950]]
```

confusion_matrix of testing set is:

```
[[146 118]
 [ 78 215]]
```

	precision	recall	f1-score	support
0	0.65	0.55	0.60	264
1	0.65	0.73	0.69	293
avg / total	0.65	0.65	0.64	557





Fold: 4

Accuracy Score: 0.6762589928057554

Precision Score: 0.6609195402298851

Recall Score: 0.7876712328767124

f1 Score: 0.71875

confusion_matrix of training set is:

```
[[612 444]
```



```
[216 955]]
```

confusion_matrix of testing set is:

```
[[146 118]
```

```
[ 62 230]]
```

	precision	recall	f1-score	support
0	0.70	0.55	0.62	264
1	0.66	0.79	0.72	292
avg / total	0.68	0.68	0.67	556



Fold: 5

Accuracy Score: 0.6312949640287769

Precision Score: 0.623229461756374

Recall Score: 0.7534246575342466

f1 Score: 0.6821705426356589


confusion_matrix of training set is:


```
[[655 401]
 [242 929]]
```

confusion_matrix of testing set is:

```
[[131 133]
 [ 72 220]]
```

	precision	recall	f1-score	support
0	0.65	0.50	0.56	264
1	0.62	0.75	0.68	292
avg / total	0.63	0.63	0.62	556





```
Out[126]:
```

	accuracy	precision	recall	f1
0	0.622980	0.617564	0.744027	0.674923
1	0.615799	0.606469	0.767918	0.677711
2	0.648115	0.645646	0.733788	0.686901
3	0.676259	0.660920	0.787671	0.718750
4	0.631295	0.623229	0.753425	0.682171

```
In [172]: clf2 = joblib.load('ClimateTeam7PD2.pkl')
```

```
In [215]: df1 = ds.Table.read_table('Climate1SupportiveLevel.csv', sep=',')
df2 = ds.Table.read_table('ClimateBalancedDS2.csv', sep=',')
df = df1.append(df2)
X = list(df['Text'])
y = list(df['Support'])
y_pred = clf_tmp.predict(X)
f1_score(y, y_pred)
```

```
Out[215]: 0.728488902401946
```